

# 01\_\_extract\_\_onsets

July 10, 2026

## 1 Onset extraction: metronome clicks + tapping onsets/accents

2026-07-10 Roberto Barumerli

Extracts, per trial: - metronome click times (from \*\_metronome.wav) - tap onset times + accent amplitude (from the plain mic \*.wav)

and saves tidy event tables + per-trial QC plots.

Run with the `bayesian_listener` conda environment (librosa installed there).

```
[1]: import re
from pathlib import Path

import librosa
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy import signal
from scipy.signal import find_peaks

RNG_SEED = 0
```

### 1.1 0. Config

The whole pipeline (Sections 1-10) runs once per subject. `SUBJECTS` is discovered from every `subNN` folder under `save/` that has a log csv, so adding a new subject's data is enough to have them picked up here without editing this notebook.

```
[2]: DATA_ROOT = Path("/Users/robaru/repos/meterswitch/save")
OUT_ROOT = Path("/Users/robaru/repos/meterswitch/save/derived")
PHASES = ["famil", "prime", "test"] # folders to scan for each subject

SUBJECTS = sorted(
    p.name for p in DATA_ROOT.iterdir()
    if p.is_dir() and (p / f"{p.name}_log.csv").exists()
)
print("Subjects found:", SUBJECTS)
```

Subjects found: ['sub01', 'sub02']

```
[3]: # --- everything below runs once per subject in SUBJECTS ---
SUBJECT = SUBJECTS[0] # set for interactive/step-by-step use; overridden by
↳ the batch loop at the end

SUBJECT_DIR = DATA_ROOT / SUBJECT
LOG_CSV = SUBJECT_DIR / f"{SUBJECT}_log.csv"

OUT_DIR = OUT_ROOT / SUBJECT
QC_DIR = OUT_DIR / "qc_plots"
OUT_DIR.mkdir(parents=True, exist_ok=True)
QC_DIR.mkdir(parents=True, exist_ok=True)

print("Subject dir:", SUBJECT_DIR)
print("Log csv:", LOG_CSV)
```

```
Subject dir: /Users/robaru/repos/meterswitch/save/sub01
Log csv: /Users/robaru/repos/meterswitch/save/sub01/sub01_log.csv
```

## 1.2 1. Build the trial manifest

Join sub01\_log.csv with the actual wav files found on disk. File naming: trial{TrialNumber}\_{meter}\_{condition}.wav and the metronome twin trial{TrialNumber}\_{meter}\_{condition}\_metronome.wav, except spontaneous test trials which carry an extra \_spontaneous / \_cued suffix before \_metronome (e.g. trial1\_3\_spontaneous.wav).

```
[4]: df_log = pd.read_csv(LOG_CSV)
df_log.columns = [c.strip() for c in df_log.columns]
df_log
```

```
[4]:
```

|    | condition   | meter | BPM | SwitchTime | Subject | TrialNumber |
|----|-------------|-------|-----|------------|---------|-------------|
| 0  | spontaneous | 3     | 120 | 0          | sub01   | 1           |
| 1  | cued        | 3     | 104 | 14         | sub01   | 2           |
| 2  | cued        | 2     | 108 | 16         | sub01   | 3           |
| 3  | spontaneous | 2     | 116 | 0          | sub01   | 4           |
| 4  | spontaneous | 3     | 108 | 0          | sub01   | 5           |
| 5  | spontaneous | 2     | 112 | 0          | sub01   | 6           |
| 6  | cued        | 3     | 112 | 12         | sub01   | 7           |
| 7  | spontaneous | 3     | 112 | 0          | sub01   | 8           |
| 8  | spontaneous | 3     | 100 | 0          | sub01   | 9           |
| 9  | cued        | 2     | 112 | 16         | sub01   | 10          |
| 10 | spontaneous | 2     | 104 | 0          | sub01   | 11          |
| 11 | spontaneous | 2     | 120 | 0          | sub01   | 12          |
| 12 | cued        | 2     | 100 | 14         | sub01   | 13          |
| 13 | cued        | 3     | 104 | 12         | sub01   | 14          |
| 14 | spontaneous | 3     | 104 | 0          | sub01   | 15          |
| 15 | cued        | 3     | 112 | 16         | sub01   | 16          |
| 16 | spontaneous | 2     | 112 | 0          | sub01   | 17          |

|    |      |   |     |    |       |    |
|----|------|---|-----|----|-------|----|
| 17 | cued | 3 | 112 | 18 | sub01 | 18 |
| 18 | cued | 2 | 100 | 18 | sub01 | 19 |
| 19 | cued | 2 | 116 | 12 | sub01 | 20 |

```
[5]: def find_trial_files(phase_dir: Path, trial_number: int, meter: int):
    """Return (mic_path, metronome_path) for a given trial in phase_dir, or
    ↪(None, None)."""
    candidates = sorted(phase_dir.glob(f"trial{trial_number}_{meter}*.wav"))
    mic_path, metro_path = None, None
    for p in candidates:
        if p.stem.endswith("_metronome"):
            if metro_path is None:
                metro_path = p
        else:
            if mic_path is None:
                mic_path = p
    return mic_path, metro_path

manifest_rows = []
for phase in PHASES:
    phase_dir = SUBJECT_DIR / phase
    if not phase_dir.exists():
        continue
    for _, row in df_log.iterrows():
        mic_path, metro_path = find_trial_files(phase_dir,
        ↪int(row["TrialNumber"]), int(row["meter"]))
        if mic_path is None and metro_path is None:
            continue # this trial doesn't live in this phase folder
        manifest_rows.append({
            "subject": SUBJECT,
            "phase": phase,
            "trial": int(row["TrialNumber"]),
            "condition": row["condition"],
            "meter": int(row["meter"]),
            "bpm": float(row["BPM"]),
            "switch_time": int(row["SwitchTime"]),
            "mic_path": mic_path,
            "metro_path": metro_path,
        })

df_manifest = pd.DataFrame(manifest_rows)
missing = df_manifest[df_manifest["mic_path"].isna() |
    ↪df_manifest["metro_path"].isna()]
if len(missing):
    print(f"WARNING: {len(missing)} trial(s) missing one of the two wav files:")
    print(missing[["phase", "trial", "condition"]])
```

```
df_manifest = df_manifest.dropna(subset=["mic_path", "metro_path"]).
↳reset_index(drop=True)
print(f"Manifest: {len(df_manifest)} trials found across phases_
↳{sorted(df_manifest['phase'].unique())}")
df_manifest.head()
```

WARNING: 21 trial(s) missing one of the two wav files:

|    | phase | trial | condition   |
|----|-------|-------|-------------|
| 0  | famil | 2     | cued        |
| 1  | prime | 1     | spontaneous |
| 2  | prime | 2     | cued        |
| 3  | prime | 3     | cued        |
| 4  | prime | 4     | spontaneous |
| 5  | prime | 5     | spontaneous |
| 6  | prime | 6     | spontaneous |
| 7  | prime | 7     | cued        |
| 8  | prime | 8     | spontaneous |
| 9  | prime | 9     | spontaneous |
| 10 | prime | 10    | cued        |
| 11 | prime | 11    | spontaneous |
| 12 | prime | 12    | spontaneous |
| 13 | prime | 13    | cued        |
| 14 | prime | 14    | cued        |
| 15 | prime | 15    | spontaneous |
| 16 | prime | 16    | cued        |
| 17 | prime | 17    | spontaneous |
| 18 | prime | 18    | cued        |
| 19 | prime | 19    | cued        |
| 20 | prime | 20    | cued        |

Manifest: 20 trials found across phases ['test']

```
[5]:  subject phase trial condition meter bpm switch_time \
0   sub01 test    1  spontaneous    3  120.0          0
1   sub01 test    2           cued    3  104.0         14
2   sub01 test    3           cued    2  108.0         16
3   sub01 test    4  spontaneous    2  116.0          0
4   sub01 test    5  spontaneous    3  108.0          0
```

```
mic_path \
0 /Users/robaru/repos/meterswitch/save/sub01/tes...
1 /Users/robaru/repos/meterswitch/save/sub01/tes...
2 /Users/robaru/repos/meterswitch/save/sub01/tes...
3 /Users/robaru/repos/meterswitch/save/sub01/tes...
4 /Users/robaru/repos/meterswitch/save/sub01/tes...
```

```
metro_path
0 /Users/robaru/repos/meterswitch/save/sub01/tes...
```

```

1 /Users/robaru/repos/meterswitch/save/sub01/tes...
2 /Users/robaru/repos/meterswitch/save/sub01/tes...
3 /Users/robaru/repos/meterswitch/save/sub01/tes...
4 /Users/robaru/repos/meterswitch/save/sub01/tes...

```

### 1.3 2. Audio loading helpers

```

[6]: def load_mic(path: Path):
      y, sr = librosa.load(str(path), sr=None, mono=True)
      return y, sr

def load_metronome(path: Path):
    y, sr = librosa.load(str(path), sr=None, mono=False)
    if y.ndim == 2:
        # verify channels are (near-)identical; fall back to channel 0
        ↪regardless
        if y.shape[0] >= 2:
            max_diff = np.max(np.abs(y[0] - y[1]))
            if max_diff > 1e-3:
                print(f"NOTE: metronome channels differ (max abs diff {max_diff:
                ↪.4f}) for {path.name}")
            y = y[0]
        return y, sr

def check_alignment(mic_y, metro_y, sr, tol_ms=50):
    dur_mic = len(mic_y) / sr
    dur_metro = len(metro_y) / sr
    diff_ms = abs(dur_mic - dur_metro) * 1000
    if diff_ms > tol_ms:
        print(f"WARNING: mic/metronome duration mismatch: {diff_ms:.1f} ms "
              f"(mic {dur_mic:.3f}s vs metronome {dur_metro:.3f}s)")
    return diff_ms

```

### 1.4 3. Metronome click detection

Clicks are strong, short, and isochronous at the trial's BPM. Peak-pick the rectified signal with a minimum distance derived from BPM, then verify IOI regularity as a QC signal.

```

[7]: def detect_clicks(y, sr, bpm, height=0.5, min_ioi_frac=0.5):
      min_distance = int(sr * (60.0 / bpm) * min_ioi_frac)
      peaks, props = find_peaks(np.abs(y), height=height,
      ↪distance=max(min_distance, 1))
      times = peaks / sr
      iois = np.diff(times)
      expected_ioi = 60.0 / bpm

```

```

ioi_std = float(np.std(iois)) if len(iois) else np.nan
ioi_mean = float(np.mean(iois)) if len(iois) else np.nan
flagged = (
    len(times) == 0
    or (not np.isnan(ioi_std) and ioi_std > 0.15 * expected_ioi)
    or (not np.isnan(ioi_mean) and abs(ioi_mean - expected_ioi) > 0.25 *
↪expected_ioi)
)
return times, {"ioi_mean": ioi_mean, "ioi_std": ioi_std, "expected_ioi":
↪expected_ioi, "n_clicks": len(times), "flagged": flagged}

```

## 1.5 4. Tap onset detection (mic channel)

The mic signal turned out to be dominated by low-frequency rumble/hum (spectral peak ~23 Hz) sitting under the much broader-band tap transients (spectral peak ~94 Hz) — librosa’s spectral-flux `onset_detect` was tripping on that rumble and massively over-counting (~190 “taps” in a trial with only ~60 real ones). Fix: high-pass filter first to remove the rumble, then peak-pick a lightly smoothed rectified envelope.

Absolute recording level varies a lot trial-to-trial (some trials are far quieter/noisier than others), so a single fixed amplitude threshold either missed almost everything on quiet trials or over-fired on loud ones. The threshold is instead set per trial from the envelope’s own noise floor: `median + k * MAD` (a robust, outlier-resistant estimate of “typical background” for that specific recording), so a peak has to stand out from *that trial’s* noise, not from a global constant.

Some trials (see QC plots) are simply low-SNR throughout with no cleanly isolated transients — detections there are lower-confidence and worth a visual spot-check rather than blind trust.

```

[8]: def detect_taps(y, sr, highpass_hz=80, smooth_ms=3, min_gap_s=0.15, mad_k=8):
    sos = signal.butter(4, highpass_hz, btype="highpass", fs=sr, output="sos")
    y_hp = signal.sosfiltfilt(sos, y)

    env = np.abs(y_hp)
    win = max(int(sr * smooth_ms / 1000), 1)
    env_smooth = np.convolve(env, np.ones(win) / win, mode="same")

    med = np.median(env_smooth)
    mad = np.median(np.abs(env_smooth - med))
    threshold = med + mad_k * mad

    peaks, props = find_peaks(env_smooth, height=threshold, distance=int(sr *
↪min_gap_s))
    times = peaks / sr

    # how far the detected peaks stand above this trial's own noise floor --
    # a trial can have a plausible *tap count* purely because the adaptive
    # threshold always finds *some* peaks, even in pure noise, so the count
    # alone can't tell a clean detection from a marginal one. This ratio can.

```

```

    snr_ratio = float(props["peak_heights"].mean() / med) if len(peaks) and med_
↪ 0 else np.nan

    return times, env_smooth, snr_ratio

def check_metronome_bleed(tap_times, click_times, tol_s=0.02):
    """Flag suspiciously high overlap between detected taps and click times
    (possible metronome bleed into the mic channel)."""
    if len(tap_times) == 0 or len(click_times) == 0:
        return 0.0
    hits = 0
    for t in tap_times:
        if np.min(np.abs(click_times - t)) < tol_s:
            hits += 1
    return hits / len(tap_times)

def is_low_confidence_trial(n_taps, expected_n_clicks, snr_ratio,
                             tap_count_ratio_bounds=(0.5, 2.0), snr_min=10.0):
    """Flag a trial as low-confidence if either:
    - the detected tap count is implausible relative to expected beats (far
      too few/many), or
    - the detected peaks don't stand out much above that trial's own noise
      floor (snr_ratio, from detect_taps) -- the adaptive MAD threshold in
      detect_taps will always find *some* peaks even in near-pure noise, so
      tap count alone can't distinguish a clean recording from a marginal
      one; snr_ratio can.
    This is a coarse, cheap sanity check meant to steer a human to the QC
    plot, not a hard exclusion criterion."""
    if expected_n_clicks == 0:
        return True
    ratio = n_taps / expected_n_clicks
    lo, hi = tap_count_ratio_bounds
    count_flag = ratio < lo or ratio > hi
    snr_flag = np.isnan(snr_ratio) or snr_ratio < snr_min
    return bool(count_flag or snr_flag)

```

## 1.6 5. Accent / amplitude extraction per tap

For each tap onset, compute peak amplitude and RMS (in dB) in a short window following the onset.

```

[9]: def tap_accents(y, sr, tap_times, window_s=0.03):
    win = int(sr * window_s)
    peak_amp, rms_db = [], []
    for t in tap_times:

```

```

start = int(t * sr)
end = min(start + win, len(y))
seg = y[start:end]
if len(seg) == 0:
    peak_amp.append(np.nan)
    rms_db.append(np.nan)
    continue
p = float(np.max(np.abs(seg)))
rms = float(np.sqrt(np.mean(seg ** 2)))
peak_amp.append(p)
rms_db.append(20 * np.log10(max(rms, 1e-8)))
return np.array(peak_amp), np.array(rms_db)

```

## 1.7 6. Per-trial QC plot

Top: mic waveform with detected tap onsets (stem height = accent). Bottom: metronome waveform with detected clicks, plus tap times overlaid (dashed) to visually check tap/beat alignment.

```

[10]: def plot_trial_qc(trial_row, mic_y, sr_mic, metro_y, sr_metro,
                        tap_times, rms_db, click_times, out_path):
    t_mic = np.arange(len(mic_y)) / sr_mic
    t_metro = np.arange(len(metro_y)) / sr_metro

    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 5), sharex=True,
    ↪ layout="constrained")

    ax1.plot(t_mic, mic_y, color="0.6", lw=0.5)
    if len(tap_times):
        norm = (rms_db - np.nanmin(rms_db)) / (np.ptp(rms_db) + 1e-9)
        ax1.vlines(tap_times, 0, norm * np.nanmax(np.abs(mic_y)), color="C3",
    ↪ lw=1.2)
        ax1.scatter(tap_times, np.full_like(tap_times, np.nanmax(np.abs(mic_y)))
    ↪ * 1.05),
                        color="C3", marker="v", s=15, label="tap onset")
    ax1.set_ylabel("mic amplitude")
    ax1.legend(fontsize=8, loc="upper right")

    ax2.plot(t_metro, metro_y, color="0.6", lw=0.5)
    if len(click_times):
        ax2.vlines(click_times, -1, 1, color="C0", lw=0.8, label="click")
    if len(tap_times):
        ax2.vlines(tap_times, -1, 1, color="C3", lw=0.8, ls="--", alpha=0.7,
    ↪ label="tap (overlay)")
    ax2.set_ylabel("metronome amplitude")
    ax2.set_xlabel("time (s)")
    ax2.legend(fontsize=8, loc="upper right")

```



```

fig.suptitle(
    f"{trial_row['subject']} trial {trial_row['trial']}\n
    ↳({trial_row['phase']}) - "
    f"{trial_row['condition']}, meter {trial_row['meter']},\n
    ↳({trial_row['bpm']:.0f} BPM, "
    f"switch_time={trial_row['switch_time']}",
    fontsize=10,
)
fig.savefig(out_path, dpi=110)
plt.close(fig)

```

## 1.8 7. Batch driver: run the pipeline over every trial in the manifest

```

[11]: metronome_events = []
tap_events = []
trial_qc_summary = []

for _, row in df_manifest.iterrows():
    mic_y, sr_mic = load_mic(row["mic_path"])
    metro_y, sr_metro = load_metronome(row["metro_path"])
    assert sr_mic == sr_metro, f"sample rate mismatch for trial {row['trial']}"
    sr = sr_mic

    dur_diff_ms = check_alignment(mic_y, metro_y, sr)

    click_times, click_qc = detect_clicks(metro_y, sr, row["bpm"])
    tap_times, onset_env, snr_ratio = detect_taps(mic_y, sr)
    peak_amp, rms_db = tap_accents(mic_y, sr, tap_times)
    bleed_frac = check_metronome_bleed(tap_times, click_times)

    for i, ct in enumerate(click_times):
        metronome_events.append({
            "subject": row["subject"], "phase": row["phase"], "trial":\n
            ↳row["trial"],
            "condition": row["condition"], "meter": row["meter"], "bpm":\n
            ↳row["bpm"],
            "switch_time": row["switch_time"], "click_idx": i, "click_time_s":\n
            ↳ct,
        })

    for i, (tt, pa, rd) in enumerate(zip(tap_times, peak_amp, rms_db)):
        tap_events.append({
            "subject": row["subject"], "phase": row["phase"], "trial":\n
            ↳row["trial"],
            "condition": row["condition"], "meter": row["meter"], "bpm":\n
            ↳row["bpm"],

```

```

        "switch_time": row["switch_time"], "tap_idx": i, "tap_time_s": tt,
        "peak_amp": pa, "rms_db": rd,
    })

    expected_n_clicks = int(np.floor((len(metro_y) / sr) / (60.0 / row["bpm"])))
    low_confidence = is_low_confidence_trial(len(tap_times), expected_n_clicks,
    ↪snr_ratio)
    trial_qc_summary.append({
        "subject": row["subject"], "phase": row["phase"], "trial": row["trial"],
        "condition": row["condition"], "n_clicks": click_qc["n_clicks"],
        "expected_n_clicks": expected_n_clicks,
        "click_ioi_std_ms": click_qc["ioi_std"] * 1000 if not np.
    ↪isnan(click_qc["ioi_std"]) else np.nan,
        "click_flagged": click_qc["flagged"],
        "n_taps": len(tap_times),
        "snr_ratio": snr_ratio,
        "bleed_frac": bleed_frac,
        "bleed_flagged": bleed_frac > 0.3,
        "low_confidence_trial": low_confidence,
        "duration_diff_ms": dur_diff_ms,
    })

    qc_path = QC_DIR / f"{row['subject']}_trial{row['trial']}_{'phase'}.
    ↪png"
    plot_trial_qc(row, mic_y, sr, metro_y, sr, tap_times, rms_db, click_times,
    ↪qc_path)

df_metronome_events = pd.DataFrame(metronome_events)
df_tap_events = pd.DataFrame(tap_events)
df_qc = pd.DataFrame(trial_qc_summary)

print(f"{len(df_manifest)} trials processed")
print(f"Flagged trials (click irregularity, bleed, or low-confidence tap count):
    ↪")
df_qc[df_qc["click_flagged"] | df_qc["bleed_flagged"] |
    ↪df_qc["low_confidence_trial"]]

```

20 trials processed

Flagged trials (click irregularity, bleed, or low-confidence tap count):

```
[11]:
```

|   | subject | phase | trial | condition   | n_clicks | expected_n_clicks | \ |
|---|---------|-------|-------|-------------|----------|-------------------|---|
| 4 | sub01   | test  | 5     | spontaneous | 54       | 54                |   |
| 5 | sub01   | test  | 6     | spontaneous | 56       | 56                |   |
| 6 | sub01   | test  | 7     | cued        | 56       | 56                |   |
| 7 | sub01   | test  | 8     | spontaneous | 56       | 56                |   |
| 8 | sub01   | test  | 9     | spontaneous | 50       | 50                |   |
| 9 | sub01   | test  | 10    | cued        | 56       | 56                |   |

|    |       |      |    |             |    |    |
|----|-------|------|----|-------------|----|----|
| 10 | sub01 | test | 11 | spontaneous | 52 | 52 |
| 11 | sub01 | test | 12 | spontaneous | 60 | 60 |
| 12 | sub01 | test | 13 | cued        | 50 | 50 |
| 13 | sub01 | test | 14 | cued        | 52 | 52 |
| 14 | sub01 | test | 15 | spontaneous | 52 | 52 |
| 15 | sub01 | test | 16 | cued        | 56 | 56 |
| 16 | sub01 | test | 17 | spontaneous | 56 | 56 |
| 17 | sub01 | test | 18 | cued        | 56 | 56 |
| 18 | sub01 | test | 19 | cued        | 50 | 50 |

|    | click_ioi_std_ms | click_flagged | n_taps | snr_ratio | bleed_frac | \ |
|----|------------------|---------------|--------|-----------|------------|---|
| 4  | 1.129330e-12     | False         | 70     | 7.561511  | 0.071429   |   |
| 5  | 1.217017e-12     | False         | 63     | 5.394053  | 0.111111   |   |
| 6  | 1.217017e-12     | False         | 65     | 6.173234  | 0.230769   |   |
| 7  | 1.217017e-12     | False         | 67     | 8.074010  | 0.149254   |   |
| 8  | 1.245617e-12     | False         | 61     | 8.191906  | 0.065574   |   |
| 9  | 1.217017e-12     | False         | 50     | 5.538535  | 0.060000   |   |
| 10 | 1.158585e-12     | False         | 49     | 4.854714  | 0.102041   |   |
| 11 | 2.653378e-13     | False         | 67     | 5.560552  | 0.089552   |   |
| 12 | 1.245617e-12     | False         | 54     | 6.518899  | 0.092593   |   |
| 13 | 1.158585e-12     | False         | 61     | 5.914323  | 0.032787   |   |
| 14 | 1.158585e-12     | False         | 61     | 8.021395  | 0.032787   |   |
| 15 | 1.217017e-12     | False         | 72     | 6.216627  | 0.069444   |   |
| 16 | 1.217017e-12     | False         | 72     | 6.805296  | 0.083333   |   |
| 17 | 1.217017e-12     | False         | 72     | 8.801102  | 0.083333   |   |
| 18 | 1.245617e-12     | False         | 64     | 6.635606  | 0.031250   |   |

|    | bleed_flagged | low_confidence_trial | duration_diff_ms |
|----|---------------|----------------------|------------------|
| 4  | False         | True                 | 0.354167         |
| 5  | False         | True                 | 0.354167         |
| 6  | False         | True                 | 28.312500        |
| 7  | False         | True                 | 0.354167         |
| 8  | False         | True                 | 0.354167         |
| 9  | False         | True                 | 28.312500        |
| 10 | False         | True                 | 0.354167         |
| 11 | False         | True                 | 0.354167         |
| 12 | False         | True                 | 28.312500        |
| 13 | False         | True                 | 28.312500        |
| 14 | False         | True                 | 0.354167         |
| 15 | False         | True                 | 28.312500        |
| 16 | False         | True                 | 0.354167         |
| 17 | False         | True                 | 28.312500        |
| 18 | False         | True                 | 28.312500        |

## 1.9 8. Save event tables + QC summary


```
[12]: df_metronome_events.to_csv(OUT_DIR / f"{SUBJECT}_metronome_events.csv",  
    ↪index=False)  
df_tap_events.to_csv(OUT_DIR / f"{SUBJECT}_tap_events.csv", index=False)  
df_qc.to_csv(OUT_DIR / f"{SUBJECT}_qc_summary.csv", index=False)  
  
print("Saved:")  
print(" -", OUT_DIR / f"{SUBJECT}_metronome_events.csv")  
print(" -", OUT_DIR / f"{SUBJECT}_tap_events.csv")  
print(" -", OUT_DIR / f"{SUBJECT}_qc_summary.csv")  
print(" - per-trial QC plots in", QC_DIR)
```

Saved:

- /Users/robaru/repos/meterswitch/derived/sub01/sub01\_metronome\_events.csv
- /Users/robaru/repos/meterswitch/derived/sub01/sub01\_tap\_events.csv
- /Users/robaru/repos/meterswitch/derived/sub01/sub01\_qc\_summary.csv
- per-trial QC plots in /Users/robaru/repos/meterswitch/derived/sub01/qc\_plots

## 1.10 9. Summary plots across all trials

```
[13]: fig, axes = plt.subplots(1, 3, figsize=(13, 3.5), layout="constrained")  
  
axes[0].scatter(df_qc["expected_n_clicks"], df_qc["n_clicks"], color="C0")  
lims = [0, max(df_qc["expected_n_clicks"].max(), df_qc["n_clicks"].max()) + 5]  
axes[0].plot(lims, lims, "k--", lw=1)  
axes[0].set_xlabel("expected n clicks (BPM x duration)")  
axes[0].set_ylabel("detected n clicks")  
axes[0].set_title("Click count check")  
  
axes[1].hist(df_tap_events["rms_db"].dropna(), bins=30, color="C3")  
axes[1].set_xlabel("tap RMS (dB)")  
axes[1].set_title("Tap accent distribution")  
  
axes[2].hist(df_qc["click_ioi_std_ms"].dropna(), bins=30, color="C0")  
axes[2].set_xlabel("click IOI std (ms)")  
axes[2].set_title("Click regularity per trial")  
  
fig.suptitle(f"{SUBJECT}: extraction summary across {len(df_manifest)} trials",  
    ↪fontsize=11)  
plt.show()
```



### 1.11 10. Align taps to clicks: per-click asynchrony, amplitude, and accent label

For each metronome click, find the closest tap within a tolerance window (a fraction of the beat period) and record the signed asynchrony `tap_time - click_time`. Once a tap is matched to a click it is “consumed” and cannot be matched to a later click — this stops one loud/ambiguous tap from being double-counted across two adjacent beats, and clicks with no available tap in range are left as missed beats (`asynchrony_ms = NaN`), rather than snapping to a distant unrelated tap.

The matched tap’s amplitude (`peak_amp`, `rms_db`, from Section 5) is carried along so we know how hard that beat was tapped. From that, each matched beat gets an `accented` label: True if its `rms_db` is above the *median rms\_db for that trial* — accent is judged relative to how the participant tapped within that specific trial, not against a fixed loudness cutoff, since absolute recording level varies a lot across trials (see Section 4). Missed beats get `accented = NA` (no tap, nothing to judge).

`low_confidence_trial` (from the QC summary, Section 7) is also carried into this table so a trial with an implausible tap count can be filtered or down-weighted directly from this one file, without a join.

This is computed post-hoc from the already-saved event tables, so it can be re-run/re-tuned (tolerance fraction, accent rule) without redoing onset detection.

```
[14]: def match_clicks_to_taps(click_times, tap_times, tap_rms_db, bpm, tol_frac=0.4):
    """For each click (in time order), claim the closest not-yet-used tap
    within tol_frac * beat_period. Returns one row per click, including the
    matched tap's rms_db and a within-trial relative accent label."""
    beat_period = 60.0 / bpm
    tol = tol_frac * beat_period

    tap_times = np.asarray(tap_times)
    tap_rms_db = np.asarray(tap_rms_db)
    used = np.zeros(len(tap_times), dtype=bool)

    rows = []
    for click_idx, ct in enumerate(click_times):
        available = ~used
        if not np.any(available):
```

```

        rows.append((click_idx, ct, np.nan, np.nan, np.nan, np.nan))
        continue
    avail_idx = np.where(available)[0]
    dists = np.abs(tap_times[avail_idx] - ct)
    best = avail_idx[np.argmin(dists)]
    best_dist = dists.min()
    if best_dist <= tol:
        used[best] = True
        asynchrony_s = tap_times[best] - ct
        rows.append((click_idx, ct, tap_times[best], int(best),
↪asynchrony_s, tap_rms_db[best]))
    else:
        rows.append((click_idx, ct, np.nan, np.nan, np.nan, np.nan))

out = pd.DataFrame(rows, columns=[
    "click_idx", "click_time_s", "matched_tap_time_s", "matched_tap_idx",
    "asynchrony_s", "rms_db",
])
out["asynchrony_ms"] = out["asynchrony_s"] * 1000
out["missed_beat"] = out["matched_tap_time_s"].isna()

median_rms_db = out["rms_db"].median() # median over matched taps in this
↪trial
out["accented"] = np.where(out["missed_beat"], np.nan, out["rms_db"] >
↪median_rms_db)
return out

```

TOL\_FRAC = 0.4 # fraction of beat period allowed as matching tolerance

```

alignment_rows = []
for _, row in df_manifest.iterrows():
    ct = df_metronome_events.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↪trial"
    )["click_time_s"].values
    tap_sub = df_tap_events.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↪trial"
    )
    tt = tap_sub["tap_time_s"].values
    tap_rms = tap_sub["rms_db"].values

    m = match_clicks_to_taps(ct, tt, tap_rms, row["bpm"], tol_frac=TOL_FRAC)
    m.insert(0, "trial", row["trial"])
    m.insert(0, "phase", row["phase"])
    m.insert(0, "subject", row["subject"])

```

```

m["condition"] = row["condition"]
m["meter"] = row["meter"]
m["bpm"] = row["bpm"]
m["switch_time"] = row["switch_time"]
low_conf = df_qc.query(
    "subject == @row.subject and phase == @row.phase and trial == @row.
↳trial"
)["low_confidence_trial"].iloc[0]
m["low_confidence_trial"] = low_conf
alignment_rows.append(m)

df_alignment = pd.concat(alignment_rows, ignore_index=True)
df_alignment.to_csv(OUT_DIR / f"{SUBJECT}_click_tap_alignment.csv", index=False)

n_missed = df_alignment["missed_beat"].sum()
n_low_conf_trials = df_qc["low_confidence_trial"].sum()
print(f"{len(df_alignment)} clicks total, {n_missed} missed beats "
      f"({100 * n_missed / len(df_alignment):.1f}%)")
print(f"{n_low_conf_trials} / {len(df_manifest)} trials flagged_
↳low_confidence_trial")
print("Saved:", OUT_DIR / f"{SUBJECT}_click_tap_alignment.csv")
df_alignment.head(10)

```

1092 clicks total, 152 missed beats (13.9%)

15 / 20 trials flagged low\_confidence\_trial

Saved:

/Users/robaru/repos/meterswitch/derived/sub01/sub01\_click\_tap\_alignment.csv

```

[14]:  subject phase  trial  click_idx  click_time_s  matched_tap_time_s  \
0    sub01  test      1           0      0.024438                NaN
1    sub01  test      1           1      0.524438             0.493521
2    sub01  test      1           2      1.024437             1.031333
3    sub01  test      1           3      1.524437             1.369208
4    sub01  test      1           4      2.024437             1.912854
5    sub01  test      1           5      2.524437             2.425500
6    sub01  test      1           6      3.024437             2.872021
7    sub01  test      1           7      3.524437             3.372500
8    sub01  test      1           8      4.024438             3.920583
9    sub01  test      1           9      4.524438                NaN

      matched_tap_idx  asynchrony_s      rms_db  asynchrony_ms  missed_beat  \
0                NaN           NaN           NaN           NaN           True
1                0.0      -0.030917 -51.116308      -30.916667          False
2                1.0       0.006896 -59.799739       6.895833          False
3                2.0      -0.155229 -60.962212     -155.229167          False
4                3.0      -0.111583 -49.567057    -111.583333          False
5                4.0      -0.098937 -55.363678     -98.937500          False

```

|   |     |           |            |             |       |
|---|-----|-----------|------------|-------------|-------|
| 6 | 5.0 | -0.152417 | -53.201601 | -152.416667 | False |
| 7 | 6.0 | -0.151937 | -52.209202 | -151.937500 | False |
| 8 | 7.0 | -0.103854 | -56.346986 | -103.854167 | False |
| 9 | NaN | NaN       | NaN        | NaN         | True  |

|   | accented | condition   | meter | bpm   | switch_time | low_confidence_trial |
|---|----------|-------------|-------|-------|-------------|----------------------|
| 0 | NaN      | spontaneous | 3     | 120.0 | 0           | False                |
| 1 | 1.0      | spontaneous | 3     | 120.0 | 0           | False                |
| 2 | 0.0      | spontaneous | 3     | 120.0 | 0           | False                |
| 3 | 0.0      | spontaneous | 3     | 120.0 | 0           | False                |
| 4 | 1.0      | spontaneous | 3     | 120.0 | 0           | False                |
| 5 | 0.0      | spontaneous | 3     | 120.0 | 0           | False                |
| 6 | 1.0      | spontaneous | 3     | 120.0 | 0           | False                |
| 7 | 1.0      | spontaneous | 3     | 120.0 | 0           | False                |
| 8 | 0.0      | spontaneous | 3     | 120.0 | 0           | False                |
| 9 | NaN      | spontaneous | 3     | 120.0 | 0           | False                |

### 1.11.1 Missed-beat rate per trial and asynchrony distribution

```
[15]: missed_per_trial = df_alignment.groupby(["trial", "condition"])["missed_beat"].
      ↪mean().reset_index()
missed_per_trial = missed_per_trial.rename(columns={"missed_beat": "missed_beat_frac"})


fig, axes = plt.subplots(1, 2, figsize=(11, 3.5), layout="constrained")

axes[0].bar(missed_per_trial["trial"].astype(str),
      ↪missed_per_trial["missed_beat_frac"] * 100,
      color=["C0" if c == "cued" else "C3" for c in
      ↪missed_per_trial["condition"]])
axes[0].set_xlabel("trial")
axes[0].set_ylabel("missed beats (%)")
axes[0].set_title("Missed-beat rate per trial")
axes[0].tick_params(axis="x", labelrotation=90, labelsize=7)

axes[1].hist(df_alignment["asynchrony_ms"].dropna(), bins=40, color="C0")
axes[1].axvline(0, color="k", ls="--", lw=1)
axes[1].set_xlabel("asynchrony (ms), tap - click")
axes[1].set_title("Tap-click asynchrony distribution")

fig.suptitle(f"{SUBJECT}: click-tap alignment (tolerance = {TOL_FRAC:.0%} of
      ↪beat period)", fontsize=11)
plt.show()
```





### 1.11.2 Per-trial alignment QC plot

Extends the Section 6 plot: matched clicks are marked, missed beats are highlighted, to visually confirm the matching makes sense against the raw waveforms.

```
[16]: def plot_trial_alignment(trial_row, mic_y, sr, click_times, tap_times,
    ↪ alignment_trial, out_path):
    t_mic = np.arange(len(mic_y)) / sr
    fig, ax = plt.subplots(figsize=(12, 3), layout="constrained")
    ax.plot(t_mic, mic_y, color="0.7", lw=0.5)

    matched = alignment_trial.dropna(subset=["matched_tap_time_s"])
    missed = alignment_trial[alignment_trial["missed_beat"]]

    ymax = np.max(np.abs(mic_y))
    ax.vlines(click_times, -ymax * 1.2, ymax * 1.2, color="C0", lw=0.7, alpha=0.
    ↪ 5, label="click")
    ax.scatter(matched["matched_tap_time_s"], np.full(len(matched), ymax * 1.1),
               color="C2", marker="v", s=25, zorder=5, label="matched tap")
    ax.vlines(missed["click_time_s"], -ymax * 1.2, ymax * 1.2, color="red",
    ↪ lw=1.5, ls=":", label="missed beat")

    ax.set_ylim(-ymax * 1.3, ymax * 1.3)
    ax.set_xlabel("time (s)")
    ax.set_ylabel("mic amplitude")
    ax.legend(fontsize=8, loc="upper right", ncol=3)
    fig.suptitle(
        f"{trial_row['subject']} trial {trial_row['trial']} - click/tap
    ↪ alignment "
        f"({alignment_trial['missed_beat'].sum()} missed beats)",
        fontsize=10,
    )
    fig.savefig(out_path, dpi=110)
    plt.close(fig)
```

```

ALIGN_QC_DIR = OUT_DIR / "alignment_qc_plots"
ALIGN_QC_DIR.mkdir(exist_ok=True)

for _, row in df_manifest.iterrows():
    mic_y, sr = load_mic(row["mic_path"])
    ct = df_metronome_events.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↳trial"
    )["click_time_s"].values
    tt = df_tap_events.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↳trial"
    )["tap_time_s"].values
    align_trial = df_alignment.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↳trial"
    )
    out_path = ALIGN_QC_DIR /
↳f"{row['subject']}_{trial}{row['trial']}_{row['phase']}_alignment.png"
    plot_trial_alignment(row, mic_y, sr, ct, tt, align_trial, out_path)

print("Saved per-trial alignment QC plots to", ALIGN_QC_DIR)

```

Saved per-trial alignment QC plots to  
/Users/robaru/repos/meterswitch/derived/sub01/alignment\_qc\_plots

## 1.12 11. Batch over every subject

Sections 1-10 above ran end-to-end for `SUBJECT = SUBJECTS[0]` as a walkthrough. `run_pipeline_for_subject` repeats that exact same pipeline (manifest -> click/tap detection -> QC plots -> click-tap alignment with amplitude + accent label) for a given subject, so it can be re-run for every subject found under `save/` without re-reading this whole notebook top to bottom per subject.

```

[17]: def run_pipeline_for_subject(subject):
    subject_dir = DATA_ROOT / subject
    log_csv = subject_dir / f"{subject}_log.csv"
    out_dir = OUT_ROOT / subject
    qc_dir = out_dir / "qc_plots"
    align_qc_dir = out_dir / "alignment_qc_plots"
    out_dir.mkdir(parents=True, exist_ok=True)
    qc_dir.mkdir(parents=True, exist_ok=True)
    align_qc_dir.mkdir(parents=True, exist_ok=True)

    df_log = pd.read_csv(log_csv)
    df_log.columns = [c.strip() for c in df_log.columns]

```

```

manifest_rows = []
for phase in PHASES:
    phase_dir = subject_dir / phase
    if not phase_dir.exists():
        continue
    for _, row in df_log.iterrows():
        mic_path, metro_path = find_trial_files(phase_dir,
↪int(row["TrialNumber"]), int(row["meter"]))
        if mic_path is None and metro_path is None:
            continue
        manifest_rows.append({
            "subject": subject, "phase": phase, "trial":
↪int(row["TrialNumber"]),
            "condition": row["condition"], "meter": int(row["meter"]),
            "bpm": float(row["BPM"]), "switch_time": int(row["SwitchTime"]),
            "mic_path": mic_path, "metro_path": metro_path,
        })
manifest = pd.DataFrame(manifest_rows)
missing = manifest[manifest["mic_path"].isna() | manifest["metro_path"].
↪isna()]
if len(missing):
    print(f"[{subject}] WARNING: {len(missing)} trial(s) missing a wav
↪file")
manifest = manifest.dropna(subset=["mic_path", "metro_path"]).
↪reset_index(drop=True)
print(f"[{subject}] manifest: {len(manifest)} trials across phases
↪{sorted(manifest['phase'].unique())}")

metronome_events, tap_events, qc_summary = [], [], []
for _, row in manifest.iterrows():
    mic_y, sr_mic = load_mic(row["mic_path"])
    metro_y, sr_metro = load_metronome(row["metro_path"])
    assert sr_mic == sr_metro, f"sample rate mismatch for trial
↪{row['trial']}"
    sr = sr_mic

    dur_diff_ms = check_alignment(mic_y, metro_y, sr)
    click_times, click_qc = detect_clicks(metro_y, sr, row["bpm"])
    tap_times, _, snr_ratio = detect_taps(mic_y, sr)
    peak_amp, rms_db = tap_accents(mic_y, sr, tap_times)
    bleed_frac = check_metronome_bleed(tap_times, click_times)

    for i, ct in enumerate(click_times):
        metronome_events.append({

```

```

        "subject": subject, "phase": row["phase"], "trial":␣
↪row["trial"],
        "condition": row["condition"], "meter": row["meter"], "bpm":␣
↪row["bpm"],
        "switch_time": row["switch_time"], "click_idx": i,␣
↪"click_time_s": ct,
    })
    for i, (tt, pa, rd) in enumerate(zip(tap_times, peak_amp, rms_db)):
        tap_events.append({
            "subject": subject, "phase": row["phase"], "trial":␣
↪row["trial"],
            "condition": row["condition"], "meter": row["meter"], "bpm":␣
↪row["bpm"],
            "switch_time": row["switch_time"], "tap_idx": i, "tap_time_s":␣
↪tt,
            "peak_amp": pa, "rms_db": rd,
        })

    expected_n_clicks = int(np.floor((len(metro_y) / sr) / (60.0 /␣
↪row["bpm"])))
    low_confidence = is_low_confidence_trial(len(tap_times),␣
↪expected_n_clicks, snr_ratio)
    qc_summary.append({
        "subject": subject, "phase": row["phase"], "trial": row["trial"],
        "condition": row["condition"], "n_clicks": click_qc["n_clicks"],
        "expected_n_clicks": expected_n_clicks,
        "click_ioi_std_ms": click_qc["ioi_std"] * 1000 if not np.
↪isnan(click_qc["ioi_std"]) else np.nan,
        "click_flagged": click_qc["flagged"], "n_taps": len(tap_times),
        "snr_ratio": snr_ratio,
        "bleed_frac": bleed_frac, "bleed_flagged": bleed_frac > 0.3,
        "low_confidence_trial": low_confidence, "duration_diff_ms":␣
↪dur_diff_ms,
    })

    qc_path = qc_dir / f"{subject}_trial{row['trial']}_{'phase'}.png"
    plot_trial_qc(row, mic_y, sr, metro_y, sr, tap_times, rms_db,␣
↪click_times, qc_path)

df_metro_ev = pd.DataFrame(metronome_events)
df_tap_ev = pd.DataFrame(tap_events)
df_qc_sub = pd.DataFrame(qc_summary)

df_metro_ev.to_csv(out_dir / f"{subject}_metronome_events.csv", index=False)
df_tap_ev.to_csv(out_dir / f"{subject}_tap_events.csv", index=False)
df_qc_sub.to_csv(out_dir / f"{subject}_qc_summary.csv", index=False)

```

```

alignment_rows = []
for _, row in manifest.iterrows():
    ct = df_metro_ev.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↪trial"
    )["click_time_s"].values
    tap_sub = df_tap_ev.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↪trial"
    )
    tt = tap_sub["tap_time_s"].values
    tap_rms = tap_sub["rms_db"].values

    m = match_clicks_to_taps(ct, tt, tap_rms, row["bpm"], tol_frac=TOL_FRAC)
    m.insert(0, "trial", row["trial"])
    m.insert(0, "phase", row["phase"])
    m.insert(0, "subject", subject)
    m["condition"] = row["condition"]
    m["meter"] = row["meter"]
    m["bpm"] = row["bpm"]
    m["switch_time"] = row["switch_time"]
    low_conf = df_qc_sub.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↪trial"
    )["low_confidence_trial"].iloc[0]
    m["low_confidence_trial"] = low_conf
    alignment_rows.append(m)

df_align_sub = pd.concat(alignment_rows, ignore_index=True)
df_align_sub.to_csv(out_dir / f"{subject}_click_tap_alignment.csv",
↪index=False)

for _, row in manifest.iterrows():
    mic_y, sr = load_mic(row["mic_path"])
    ct = df_metro_ev.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↪trial"
    )["click_time_s"].values
    tt = df_tap_ev.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↪trial"
    )["tap_time_s"].values
    align_trial = df_align_sub.query(
        "subject == @row.subject and phase == @row.phase and trial == @row.
↪trial"
    )

```

```

    )
    out_path = align_qc_dir / "\
↪f"{subject}_trial{row['trial']}_row['phase']}_alignment.png"
    plot_trial_alignment(row, mic_y, sr, ct, tt, align_trial, out_path)

    n_missed = df_align_sub["missed_beat"].sum()
    n_low_conf = df_qc_sub["low_confidence_trial"].sum()
    print(f"[{subject}] {len(df_align_sub)} clicks, {n_missed} missed beats "
          f"({100 * n_missed / len(df_align_sub):.1f}%), "
          f"{n_low_conf}/{len(manifest)} low-confidence trials")

    return df_metro_ev, df_tap_ev, df_qc_sub, df_align_sub

```

```

[18]: # run for every subject discovered in Section 0 (SUBJECTS[0] was already run
# above as the interactive walkthrough; re-running it here is cheap and
# keeps this loop simple and uniform across all subjects)
results_by_subject = {}
for subject in SUBJECTS:
    results_by_subject[subject] = run_pipeline_for_subject(subject)

print("\nDone. Subjects processed:", list(results_by_subject.keys()))

```

```

[sub01] WARNING: 21 trial(s) missing a wav file
[sub01] manifest: 20 trials across phases ['test']

[sub01] 1092 clicks, 152 missed beats (13.9%), 15/20 low-confidence trials
[sub02] WARNING: 20 trial(s) missing a wav file
[sub02] manifest: 20 trials across phases ['test']

[sub02] 1069 clicks, 12 missed beats (1.1%), 0/20 low-confidence trials

Done. Subjects processed: ['sub01', 'sub02']

```