

02_fit_model

July 10, 2026

1 Fitting the barebone soft-change-point model to sub02's real tap data

2026-07-10 Roberto Barumerli

Fits model per-trial to real data from `save/derived/sub02/sub02_click_tap_alignment.csv`

Two stages, each trial fit independently (no cross-trial hierarchy), both using the same broad, uninformative prior on `tau` (see `build_cued_trial_model`'s docstring for why the cued model no longer gets a tight cue-anchored prior – it was masking whether the likelihood itself works): 1. **Cued** — recovers a switch latency to compare against the experimenter's `switch_time`, as a check on the model. 2. **Spontaneous** — the posterior over `tau` *is* the estimate of when the participant decided to switch on their own. Always assumes a switch happens somewhere in the trial (no lapse mixture – see `build_spontaneous_trial_model`'s docstring for why that was dropped); a trial with no clear switch shows up as a wide, unidentified `tau` posterior instead.

Model structure (revised): `meter` in the trial log is the trial's **starting** meter – participants are tasked with switching to the *other* meter (2<->3). Checked directly against the data (both conditions): grouping taps by the starting meter's period gives high accent contrast before the switch and near- zero after; grouping the same taps by the *other* meter's period gives near- zero contrast before and high contrast after – a clean swap between two periodic accent patterns, confirmed at the known `switch_beat` in every single cued trial. An earlier version of this model used one template that faded to flat, sized to the trial's single `meter` value – that missed this swap structure entirely (verified: on spontaneous trials it called “unclear,” the other-meter accent was often just as strong as the starting-meter accent was ambiguous, e.g. trial 13's other-meter contrast rises to 8.3 dB post-switch even though the model's flat-attenuation view saw nothing to detect).

The model now mixes two templates – one sized to the starting meter, one to the other meter – via the same switch weight `w_n`, matching the original barebone/full model's `g_A/g_B` mixture, but assigning “start” and “other” per-trial from the log rather than assuming a fixed pair:

```
mu_n = (1 - w_n) * g_start[n mod start_meter] + w_n * g_other[n mod other_meter]
```

Run with the `bayesian_listener` conda environment (PyMC, ArviZ already used there for the barebone notebook).

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
```

```

import pymc as pm
import pytensor.tensor as pt
import arviz as az

az.style.use("arviz-darkgrid")
mpl.rcParams["figure.dpi"] = 110
mpl.rcParams["axes.titlesize"] = 11
mpl.rcParams["axes.labelsize"] = 10
mpl.rcParams["figure.constrained_layout.use"] = True

print("PyMC", pm.__version__, "| ArviZ", az.__version__)

```

PyMC 5.28.5 | ArviZ 0.23.4

1.1 0. Config

```

[2]: SUBJECT = "sub02"
ALIGNMENT_CSV = f"/Users/robaru/repos/meterswitch/save/derived/{SUBJECT}/
↳{SUBJECT}_click_tap_alignment.csv"
OUT_DIR = f"/Users/robaru/repos/meterswitch/save/derived/{SUBJECT}"

DB_PER_NAT = 20 / np.log(10) # rms_db = 20*log10(rms) = 8.686 * ln(rms)

```

1.2 1. Load and prepare data (Step 0 of the plan)

Drop missed beats (no observation to fit), convert `rms_db` to natural-log amplitude (`log_v`) to reuse the barebone notebook's prior scales, and convert `switch_time` (seconds) to a pulse/beat index (`switch_beat`).

```

[3]: df_raw = pd.read_csv(ALIGNMENT_CSV)
df_raw["missed_beat"] = df_raw["missed_beat"].astype(str).lower().eq("true")

df = df_raw[~df_raw["missed_beat"]].copy()
df["log_v"] = df["rms_db"] / DB_PER_NAT
df["switch_beat"] = df["switch_time"] * df["bpm"] / 60.0

print(f"{len(df_raw)} clicks total, {len(df)} with a matched tap (observation)")
print(df.groupby("condition")["trial"].nunique())

df_cued = df[df["condition"] == "cued"].copy()
df_spontaneous = df[df["condition"] == "spontaneous"].copy()

trials_cued = df_cued[["trial", "meter", "bpm", "switch_time", "switch_beat"]].
↳drop_duplicates().sort_values("trial")
trials_spontaneous = df_spontaneous[["trial", "meter", "bpm"]].
↳drop_duplicates().sort_values("trial")
trials_cued

```

```

1069 clicks total, 1057 with a matched tap (observation)
condition
cued          10
spontaneous   10
Name: trial, dtype: int64

```

```

[3]:      trial  meter    bpm  switch_time  switch_beat
112      3      3  108.0          14    25.200000
220      5      2  104.0          12    20.800000
274      6      3  100.0          18    30.000000
493     10      3  112.0          12    22.400000
601     12      2  100.0          14    23.333333
705     14      3  104.0          16    27.733333
759     15      2  112.0          12    22.400000
863     17      3  112.0          12    22.400000
969     19      2  100.0          16    26.666667
1019    20      2  100.0          18    30.000000

```

1.3 2. Shared model core: single-template attenuation

$w_n = \text{logistic}((n - \text{tau}) / s)$ as in the barebone model. Instead of mixing two templates, one fixed template g_{meter} (one offset per bar position, sum-to-zero start-meter and other-meter template swap, mixed by the same switch weight:

$$\mu_n = (1 - w_n) * g_{\text{start}}[n \bmod \text{start_meter}] + w_n * g_{\text{other}}[n \bmod \text{other_meter}]$$

σ_n keeps the same switch-related noise inflation as the barebone model. (Asynchrony was checked as a second observation channel via two hypotheses – a periodic start/other template swap, and switch-related noise inflation – and rejected: the periodic pattern was weak/inconsistent across trials, and a BADS-optimized, BIC-penalized comparison rejected noise inflation in 10/10 cued trials (mean delta-BIC = +8.8 in favour of constant noise). Loudness alone is what the model fits.)

```

[4]: def other_meter(meter):
      return 3 if meter == 2 else 2

def build_cued_trial_model(n, log_v, start_meter, switch_beat, n_pulses,
    ↪tau_prior_sd=None):
    """Single cued trial, mixing a template sized to the trial's starting meter
    with one sized to the other meter.

    tau's prior is broad/uninformative by default (same scale for both cued
    and spontaneous fits:  $n_{\text{pulses}}/2 \pm n_{\text{pulses}}/4$ ), NOT anchored tightly to
    switch_beat. This makes the cued fit a fair test of whether the
    (corrected) likelihood alone recovers the true switch time from the
    accent-template signal -- a tight cue-anchored prior would recover tau
    near switch_beat almost by construction regardless of whether the
    likelihood is doing any real work, which is what happened before the

```

missing-baseline bug was found (sigma_0 had inflated to ~4.0, an order of magnitude too large, silently swamping the real accent signal -- yet tau still looked "recovered" because the tight prior was carrying it).

sigma = n_pulses/4 (not /3): a trial with weak/ambiguous pre-switch signal (e.g. trial 12, where the starting- and other-meter contrasts were nearly tied before the true switch) let NUTS wander into a spurious second mode at tau near 0 -- a switch "at the very first pulse" is not a plausible reading of a cued paradigm where the participant starts in the cued starting meter. Tightening the prior modestly turns down the density available at the trial edges without re-introducing a tight cue anchor; tau=0 or tau=n_pulses are now several SDs out under any n_pulses in this dataset (50-60), rather than only ~1.5 SDs out under the old /3 scale."""

```
end_meter = other_meter(start_meter)
pos_start = np.mod(n, start_meter).astype(int)
pos_other = np.mod(n, end_meter).astype(int)
```

```
with pm.Model() as model:
    if tau_prior_sd is None:
        tau = pm.Normal("tau", mu=n_pulses / 2, sigma=n_pulses / 4)
    else:
        tau = pm.Normal("tau", mu=switch_beat, sigma=tau_prior_sd)
```

```
# A direct BADS/BIC comparison (soft sigmoid vs. a genuinely
# instantaneous hard-step model, same data, same templates) found the
# hard step wins in 18/20 trials -- freed of any prior, the sigmoid's
# own best-fit s repeatedly collapsed toward its lower bound (many
# trials fit s in the 0.10-0.25 range). The prior below is shifted to
# centre on that sharp scale (not fixed/hard-coded -- still a genuine,
# weakly-informative prior tau can be pulled away from by the data),
# rather than the old Normal(log(1.0), 0.5), which was centred an
# order of magnitude too wide relative to what the data supports.
log_s = pm.Normal("log_s", mu=np.log(0.15), sigma=0.4)
s = pm.Deterministic("s", pt.exp(log_s))
```

```
# log_v is centred around -5.7 (rms_db around -45 to -60 dB), not 0 --
# g_start/g_other are zero-mean by construction (sum-to-zero contrasts),
# so without this shared baseline the model has no way to match that
# offset except by inflating sigma_0 to absorb it, which was silently
# swamping the real ~0.5-scale accent signal (sigma_0 was fitting to
# ~4.0, an order of magnitude too large) and made tau/pi meaningless.
baseline = pm.Normal("baseline", mu=log_v.mean(), sigma=2.0)
```

```
g_start_raw = pm.Normal("g_start_raw", 0.0, 1.0, shape=start_meter)
g_start = pm.Deterministic("g_start", g_start_raw - g_start_raw.mean())
g_other_raw = pm.Normal("g_other_raw", 0.0, 1.0, shape=end_meter)
g_other = pm.Deterministic("g_other", g_other_raw - g_other_raw.mean())
```

```

sigma_0 = pm.HalfNormal("sigma_0", sigma=0.5)
sigma_sw = pm.HalfNormal("sigma_sw", sigma=0.5)

w = pm.math.sigmoid((n - tau) / s)
mu = baseline + (1 - w) * g_start[pos_start] + w * g_other[pos_other]
sigma = pt.sqrt(sigma_0**2 + sigma_sw**2 * 4 * w * (1 - w))

pm.Normal("log_v_obs", mu=mu, sigma=sigma, observed=log_v)

return model

```

`def build_spontaneous_trial_model(n, log_v, start_meter, n_pulses):`

"""Single spontaneous trial: broad prior on tau (no cue), assuming a switch always happens somewhere in the trial (no lapse mixture).

An earlier version added a discrete no-switch lapse mixture (logsumexp over a "switch happened" vs "stayed in the starting template the whole trial" branch, weighted by a fitted pi). Dropped: it produced real convergence failures on the harder trials (R-hat 1.2-1.5) -- a known failure mode of logsumexp mixtures when both branches explain the data comparably well, the sampler can get stuck unevenly split across chains near the switch/no-switch boundary. Per your steer, we now just assume a switch always happens and let tau's own posterior width communicate uncertainty (a very flat/uninformative tau posterior is itself evidence of "no clear switch", without needing a separate discrete pi to say so).

Structurally identical to build_cued_trial_model with a broad tau prior; kept as a separate function for a distinct call site/name in the driver code below."""

```

return build_cued_trial_model(n, log_v, start_meter, switch_beat=None,
↪n_pulses=n_pulses,
                                tau_prior_sd=None)

```

1.4 3. Cued fit — one trial at a time (Step 1)

tau's prior here is broad/uninformative (not anchored to `switch_beat`), so recovering the true switch time is a genuine test of the likelihood, not the prior doing the work – see `build_cued_trial_model`'s docstring.

```

[5]: cued_results = []
      cued_idatas = {}

      for _, row in trials_cued.iterrows():
          trial = int(row["trial"])
          sub = df_cued[df_cued["trial"] == trial]

```

```

n = sub["click_idx"].values
log_v = sub["log_v"].values
meter = int(row["meter"])
switch_beat = row["switch_beat"]
n_pulses = int(n.max()) + 1

model = build_cued_trial_model(n, log_v, meter, switch_beat, n_pulses)
with model:
    idata = pm.sample(1000, tune=1500, chains=4, cores=4, target_accept=0.9,
                      random_seed=20260703, progressbar=False,
↪init="adapt_diag")
    cued_idatas[trial] = idata

    tau_samples = idata.posterior["tau"].values.flatten()
    rhat = az.rhat(idata, var_names=["tau", "s", "sigma_0", "sigma_sw"])
    max_rhat = float(max(rhat[v].values for v in ["tau", "s", "sigma_0",
↪"sigma_sw"]))

    cued_results.append({
        "trial": trial, "meter": meter, "switch_beat": switch_beat,
        "tau_mean": tau_samples.mean(),
        "tau_lo": np.percentile(tau_samples, 3), "tau_hi": np.
↪percentile(tau_samples, 97),
        "max_rhat": max_rhat, "converged": max_rhat < 1.05,
    })
    print(f"trial {trial:2d} (meter {meter}): tau = {tau_samples.mean():.2f} "
          f"[{np.percentile(tau_samples,3):.2f}, {np.percentile(tau_samples,97):
↪.2f}] "
          f"vs switch_beat={switch_beat:.2f} max_rhat={max_rhat:.3f}"
          f"{' *** NOT CONVERGED ***' if max_rhat >= 1.05 else ''}")

df_cued_results = pd.DataFrame(cued_results)
df_cued_results

```

Initializing NUTS using adapt_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 6 seconds.

Initializing NUTS using adapt_diag...

trial 3 (meter 3): tau = 25.97 [24.08, 28.02] vs switch_beat=25.20
max_rhat=1.001

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 5 seconds.

The rhat statistic is larger than 1.01 for some parameters. This indicates problems during sampling. See <https://arxiv.org/abs/1903.08008> for details

The effective sample size per chain is smaller than 100 for some parameters. A higher number is needed for reliable rhat and ess computation. See <https://arxiv.org/abs/1903.08008> for details

Initializing NUTS using adapt_diag...

trial 5 (meter 2): tau = 24.76 [21.07, 27.24] vs switch_beat=20.80
max_rhat=1.016

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 5 seconds.

Initializing NUTS using adapt_diag...

trial 6 (meter 3): tau = 31.42 [29.91, 33.67] vs switch_beat=30.00
max_rhat=1.002

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 6 seconds.

Initializing NUTS using adapt_diag...

trial 10 (meter 3): tau = 22.42 [20.87, 24.68] vs switch_beat=22.40
max_rhat=1.002

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 4 seconds.

The rhat statistic is larger than 1.01 for some parameters. This indicates problems during sampling. See <https://arxiv.org/abs/1903.08008> for details

The effective sample size per chain is smaller than 100 for some parameters. A higher number is needed for reliable rhat and ess computation. See <https://arxiv.org/abs/1903.08008> for details

Initializing NUTS using adapt_diag...

trial 12 (meter 2): tau = 21.64 [0.00, 29.20] vs switch_beat=23.33
max_rhat=1.298 *** NOT CONVERGED ***

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
 Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 4 seconds.
 Initializing NUTS using adapt_diag..
 trial 14 (meter 3): tau = 30.40 [27.24, 32.99] vs switch_beat=27.73
 max_rhat=1.006
 Multiprocess sampling (4 chains in 4 jobs)
 NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
 Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 5 seconds.
 Initializing NUTS using adapt_diag..
 trial 15 (meter 2): tau = 24.49 [22.93, 26.74] vs switch_beat=22.40
 max_rhat=1.006
 Multiprocess sampling (4 chains in 4 jobs)
 NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
 Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 6 seconds.
 The rhat statistic is larger than 1.01 for some parameters. This indicates problems during sampling. See <https://arxiv.org/abs/1903.08008> for details
 The effective sample size per chain is smaller than 100 for some parameters. A higher number is needed for reliable rhat and ess computation. See <https://arxiv.org/abs/1903.08008> for details
 trial 17 (meter 3): tau = 25.50 [21.59, 27.52] vs switch_beat=22.40
 max_rhat=1.019
 Initializing NUTS using adapt_diag..
 Multiprocess sampling (4 chains in 4 jobs)
 NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
 Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 3 seconds.
 Initializing NUTS using adapt_diag..
 trial 19 (meter 2): tau = 29.43 [27.17, 31.52] vs switch_beat=26.67
 max_rhat=1.002
 Multiprocess sampling (4 chains in 4 jobs)
 NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
 Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 4 seconds.

trial 20 (meter 2): tau = 33.49 [31.14, 35.43] vs switch_beat=30.00
max_rhat=1.003

```
[5]:
```

	trial	meter	switch_beat	tau_mean	tau_lo	tau_hi	max_rhat	\
0	3	3	25.200000	25.974759	24.077614	28.023848	1.001194	
1	5	2	20.800000	24.755042	21.071625	27.235627	1.015730	
2	6	3	30.000000	31.418276	29.913758	33.673041	1.002480	
3	10	3	22.400000	22.422235	20.874127	24.677741	1.002060	
4	12	2	23.333333	21.638990	0.003991	29.204093	1.298277	
5	14	3	27.733333	30.397113	27.237876	32.989144	1.006299	
6	15	2	22.400000	24.490252	22.929086	26.741122	1.006306	
7	17	3	22.400000	25.502649	21.590933	27.521473	1.018818	
8	19	2	26.666667	29.428678	27.168879	31.516919	1.002313	
9	20	2	30.000000	33.492296	31.137522	35.432428	1.003327	

```

converged
0      True
1      True
2      True
3      True
4     False
5      True
6      True
7      True
8      True
9      True

```

1.5 4. Cued recovery check: posterior tau vs. experimenter's switch_beat

```
[6]: # Exclude trials that failed to converge (max_rhat >= 1.05) from the recovery
# check -- their posteriors aren't trustworthy point estimates (e.g. trial 12
# repeatedly shows this: its pre-switch accent contrast is unusually weak/
# ambiguous between the two candidate meters, which lets the sampler wander
# into a spurious mode near tau=0 rather than a genuine identifiability
# problem with the model itself).
converged = df_cued_results[df_cued_results["converged"]]
n_excluded = len(df_cued_results) - len(converged)
if n_excluded:
    excluded_trials = df_cued_results.loc[~df_cued_results["converged"],
    ↪ "trial"].tolist()
    print(f"Excluding {n_excluded} non-converged trial(s) from recovery check:
    ↪ {excluded_trials}")

fig, ax = plt.subplots(figsize=(7, 5), layout="constrained")
y = np.arange(len(df_cued_results))
is_converged = df_cued_results["converged"].values
```

```

for mask, color, label in [(is_converged, "C0", "posterior tau, converged"),
                           (~is_converged, "C3", "posterior tau, NOT_
↳converged")]:
    if mask.sum() == 0:
        continue
    ax.errorbar(df_cued_results["tau_mean"][mask], y[mask],
                xerr=[df_cued_results["tau_mean"][mask] -
↳df_cued_results["tau_lo"][mask],
                    df_cued_results["tau_hi"][mask] -
↳df_cued_results["tau_mean"][mask]],
                fmt="o", color=color, capsize=3, label=label)

ax.scatter(df_cued_results["switch_beat"], y, color="k", marker="x", s=60,
↳zorder=5,
            label="experimenter switch_beat")

ax.set_yticks(y)
ax.set_yticklabels([f"trial {t} (m{m})" for t, m in
↳zip(df_cued_results["trial"], df_cued_results["meter"])]])
ax.set_xlabel("tau (beats)")
ax.set_title(f"{SUBJECT}: cued fit, posterior tau vs. cue time (switch_beat)")
ax.legend(fontsize=8)
plt.show()

# NOTE: switch_beat is when the *cue* was given, not when the participant
# actually changed their tapping -- tau is fit purely from the accent-template
# signal in the taps themselves, so tau > switch_beat is not "error," it's the
# model correctly measuring the participant's reaction/adoption lag to the
# cue. Report this lag explicitly rather than treating switch_beat as ground
# truth for tau and folding the gap into an accuracy metric like MAE.
corr = np.corrcoef(converged["tau_mean"], converged["switch_beat"])[0, 1]
lag = converged["tau_mean"] - converged["switch_beat"]
print(f"corr(tau_mean, switch_beat) = {corr:.3f} (converged trials only,
↳n={len(converged)})")
print(f"estimated reaction/adoption lag (tau - switch_beat): "
      f"mean={lag.mean():.2f} beats, sd={lag.std():.2f}, range=[{lag.min():.
↳2f}, {lag.max():.2f}]")


```

Excluding 1 non-converged trial(s) from recovery check: [12]

```

/opt/homebrew/Caskroom/miniconda/base/envs/bayesian_listener/lib/python3.13/site
-packages/matplotlib/cbook.py:1719: FutureWarning: Calling float on a single
element Series is deprecated and will raise a TypeError in the future. Use
float(ser.iloc[0]) instead
    return math.isfinite(val)

```



```
corr(tau_mean, switch_beat) = 0.938 (converged trials only, n=9)
estimated reaction/adoption lag (tau - switch_beat): mean=2.25 beats, sd=1.30,
range=[0.02, 3.96]
```

1.6 5. Spontaneous fit — one trial at a time (Step 2)

Broad prior on `tau`, always assuming a switch happens somewhere in the trial (no lapse mixture – see `build_spontaneous_trial_model`'s docstring for why it was dropped: it caused real sampler convergence failures on the harder trials). The posterior `tau` here is the model's estimate of when the participant switched; its **credible-interval width** is now the confidence signal to look at – a trial with no clear switch in the data should show a wide, unidentified `tau` posterior rather than a wrong `pi`.

```
[7]: spontaneous_results = []
     spontaneous_idatas = {}

     for _, row in trials_spontaneous.iterrows():
         trial = int(row["trial"])
         sub = df_spontaneous[df_spontaneous["trial"] == trial]
         n = sub["click_idx"].values
         log_v = sub["log_v"].values
         meter = int(row["meter"])
```

```

n_pulses = int(n.max()) + 1

model = build_spontaneous_trial_model(n, log_v, meter, n_pulses)
with model:
    idata = pm.sample(1000, tune=1500, chains=4, cores=4, target_accept=0.9,
                      random_seed=20260703, progressbar=False,
↳init="adapt_diag")
    spontaneous_idatas[trial] = idata

    tau_samples = idata.posterior["tau"].values.flatten()
    rhat = az.rhat(idata, var_names=["tau", "s", "sigma_0", "sigma_sw"])
    max_rhat = float(max(rhat[v].values for v in ["tau", "s", "sigma_0",
↳"sigma_sw"]))

    bpm = row["bpm"]
    tau_lo, tau_hi = np.percentile(tau_samples, [3, 97])
    spontaneous_results.append({
        "trial": trial, "meter": meter, "n_pulses": n_pulses, "bpm": bpm,
        "tau_mean": tau_samples.mean(),
        "tau_lo": tau_lo, "tau_hi": tau_hi, "tau_ci_width": tau_hi - tau_lo,
        "tau_mean_s": tau_samples.mean() * 60 / bpm,
        "max_rhat": max_rhat, "converged": max_rhat < 1.05,
    })
    print(f"trial {trial:2d} (meter {meter}): tau = {tau_samples.mean():.2f} "
          f"[{tau_lo:.2f}, {tau_hi:.2f}] beats (width={tau_hi-tau_lo:.1f}) "
↳max_rhat={max_rhat:.3f}"
          f"{' *** NOT CONVERGED ***' if max_rhat >= 1.05 else ''}")

df_spontaneous_results = pd.DataFrame(spontaneous_results)
df_spontaneous_results

```

Initializing NUTS using adapt_diag..

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 5 seconds.

trial 1 (meter 2): tau = 18.99 [17.68, 21.99] beats (width=4.3) max_rhat=1.004

Initializing NUTS using adapt_diag..

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 5 seconds.

Initializing NUTS using adapt_diag..

```

trial 2 (meter 3): tau = 18.57 [15.93, 21.15] beats (width=5.2) max_rhat=1.005
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws
total) took 5 seconds.
Initializing NUTS using adapt_diag...

trial 4 (meter 3): tau = 21.11 [18.48, 22.99] beats (width=4.5) max_rhat=1.004
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws
total) took 5 seconds.
Initializing NUTS using adapt_diag...

trial 7 (meter 2): tau = 21.38 [19.40, 23.49] beats (width=4.1) max_rhat=1.002
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws
total) took 4 seconds.
Initializing NUTS using adapt_diag...

trial 8 (meter 2): tau = 23.49 [22.06, 26.20] beats (width=4.1) max_rhat=1.002
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws
total) took 5 seconds.
Initializing NUTS using adapt_diag...

trial 9 (meter 2): tau = 22.43 [18.80, 27.46] beats (width=8.7) max_rhat=1.006
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws
total) took 5 seconds.
Initializing NUTS using adapt_diag...

trial 11 (meter 3): tau = 31.77 [29.57, 33.93] beats (width=4.4) max_rhat=1.002
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

```

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 5 seconds.

Initializing NUTS using adapt_diag...

trial 13 (meter 2): tau = 22.09 [17.16, 29.38] beats (width=12.2)
max_rhat=1.007

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 4 seconds.

Initializing NUTS using adapt_diag...

trial 16 (meter 3): tau = 24.85 [23.03, 27.88] beats (width=4.8) max_rhat=1.003

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [tau, log_s, baseline, g_start_raw, g_other_raw, sigma_0, sigma_sw]

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 5 seconds.

trial 18 (meter 3): tau = 23.69 [20.98, 26.60] beats (width=5.6) max_rhat=1.002

```
[7]: trial meter n_pulses bpm tau_mean tau_lo tau_hi \
0      1      2         60 120.0 18.986476 17.682850 21.989312
1      2      3         52 104.0 18.569268 15.925999 21.148417
2      4      3         54 108.0 21.112468 18.484563 22.986649
3      7      2         57 116.0 21.380076 19.399033 23.492650
4      8      2         54 108.0 23.485493 22.058787 26.203610
5      9      2         60 120.0 22.432748 18.803978 27.464169
6     11      3         52 104.0 31.769768 29.573021 33.931539
7     13      2         54 108.0 22.085528 17.155118 29.378073
8     16      3         50 100.0 24.848982 23.028191 27.877072
9     18      3         50 100.0 23.685433 20.980874 26.596162
```

	tau_ci_width	tau_mean_s	max_rhat	converged
0	4.306462	9.493238	1.003715	True
1	5.222417	10.713039	1.004616	True
2	4.502085	11.729149	1.003911	True
3	4.093617	11.058660	1.001984	True
4	4.144823	13.047496	1.002432	True
5	8.660192	11.216374	1.005931	True
6	4.358518	18.328712	1.001627	True
7	12.222955	12.269738	1.007311	True
8	4.848881	14.909389	1.002615	True
9	5.615288	14.211260	1.001563	True


1.7 6. Spontaneous switch-time estimates

```
[8]: fig, axes = plt.subplots(1, 2, figsize=(11, 4.5), layout="constrained")

y = np.arange(len(df_spontaneous_results))
axes[0].errorbar(df_spontaneous_results["tau_mean"], y,
                 xerr=[df_spontaneous_results["tau_mean"] -
                     df_spontaneous_results["tau_lo"],
                     df_spontaneous_results["tau_hi"] -
                     df_spontaneous_results["tau_mean"]],
                 fmt="o", color="C2", capsize=3)
axes[0].set_yticks(y)
axes[0].set_yticklabels([f"trial {t} (m{m})" for t, m in
                         zip(df_spontaneous_results["trial"],
                             df_spontaneous_results["meter"])]])
axes[0].set_xlabel("tau (beats, estimated self-switch time)")
axes[0].set_title("Spontaneous: estimated switch time")

axes[1].bar([f"t{t}" for t in df_spontaneous_results["trial"]],
            df_spontaneous_results["tau_ci_width"],
            color="C3")
axes[1].set_ylabel("tau 94% CI width (beats)")
axes[1].set_title("Switch-time uncertainty per trial\n(wide = no clear switch_\nfound)")
axes[1].tick_params(axis="x", labelrotation=90, labelsz=7)

fig.suptitle(f"{SUBJECT}: spontaneous fit", fontsize=11)
plt.show()
```



1.8 7. Cued vs. spontaneous, meter 2 vs. meter 3: parameter comparisons

Rather than comparing posterior *means* across conditions/meters (which ignores each trial's own uncertainty), this pools full posterior draws from every converged trial and computes a proper 94%-HDI credible interval on each group difference, via a paired Monte Carlo resample: for each of many iterations, draw one posterior sample per trial in the group, average across trials to get one plausible value of “the group’s mean”, repeat for the other group, and take the difference. Doing this many times builds up the full distribution of the group-mean difference, not just a point estimate. A difference is only reported as credible if its 94% HDI excludes zero.

```
[9]: def extract_posterior_draws(idatas_dict, meters_dict, condition,
    ↪ var_names=("tau", "s", "sigma_sw", "sigma_0")):
    """Long-format: one row per (trial, draw), tagging condition/meter/converged
    so draws can be grouped and resampled later."""
    rows = []
    for trial, idata in idatas_dict.items():
        rhat = az.rhat(idata, var_names=["tau", "s", "sigma_0", "sigma_sw"])
        max_rhat = float(max(rhat[v].values for v in ["tau", "s", "sigma_0",
    ↪ "sigma_sw"]))
        draws = {v: idata.posterior[v].values.flatten() for v in var_names}
        n_draws = len(draws[var_names[0]])
        row = {"trial": trial, "condition": condition, "meter":
    ↪ meters_dict[trial],
                "draw": np.arange(n_draws), "max_rhat": max_rhat, "converged":
    ↪ max_rhat < 1.05}
        row.update(draws)
        rows.append(pd.DataFrame(row))
    return pd.concat(rows, ignore_index=True)

cued_meters = dict(zip(trials_cued["trial"], trials_cued["meter"]))
spontaneous_meters = dict(zip(trials_spontaneous["trial"],
    ↪ trials_spontaneous["meter"]))

draws_cued = extract_posterior_draws(cued_idatas, cued_meters, "cued")
draws_spontaneous = extract_posterior_draws(spontaneous_idatas,
    ↪ spontaneous_meters, "spontaneous")
df_posterior_draws = pd.concat([draws_cued, draws_spontaneous],
    ↪ ignore_index=True)

n_excluded_draws = df_posterior_draws.loc[~df_posterior_draws["converged"],
    ↪ "trial"].unique()
print(f"{df_posterior_draws['trial'].nunique()} trials, "
    f"{len(n_excluded_draws)} excluded from comparisons below (non-converged):
    ↪ "
    f"{sorted(n_excluded_draws.tolist())}")
```

20 trials, 1 excluded from comparisons below (non-converged): [12]


```

[10]: def hdi(samples, cred=0.94):
    samples = np.sort(samples)
    n = len(samples)
    interval_idx = int(np.floor(cred * n))
    n_intervals = n - interval_idx
    widths = samples[interval_idx:] - samples[:n_intervals]
    min_idx = np.argmin(widths)
    return samples[min_idx], samples[min_idx + interval_idx]

def group_mean_mc(df, group_col, group_val, param, n_mc=4000, rng=None):
    """Paired Monte Carlo resample of the group's mean, respecting each
    trial's own posterior uncertainty (see markdown above)."""
    rng = rng if rng is not None else np.random.default_rng(0)
    sub = df[df[group_col] == group_val]
    trial_draws = [g[param].values for _, g in sub.groupby("trial")]
    mc_means = np.empty(n_mc)
    for i in range(n_mc):
        mc_means[i] = np.mean([rng.choice(td) for td in trial_draws])
    return mc_means

def compare_groups(df, group_col, group_a, group_b, params, rng):
    rows = []
    for param in params:
        a_mc = group_mean_mc(df, group_col, group_a, param, rng=rng)
        b_mc = group_mean_mc(df, group_col, group_b, param, rng=rng)
        diff_mc = a_mc - b_mc
        lo, hi = hdi(diff_mc)
        rows.append({
            "comparison": f"{group_a} - {group_b}", "parameter": param,
            f"{group_a}_mean": a_mc.mean(), f"{group_b}_mean": b_mc.mean(),
            "diff_mean": diff_mc.mean(), "hdi_94_lo": lo, "hdi_94_hi": hi,
            "credible": not (lo <= 0 <= hi),
        })
    return pd.DataFrame(rows)

df_converged_draws = df_posterior_draws[df_posterior_draws["converged"]]
rng = np.random.default_rng(42)
params_to_check = ["tau", "s", "sigma_sw", "sigma_0"]

condition_comparison = compare_groups(df_converged_draws, "condition", "cued",
    ↪ "spontaneous",
    params_to_check, rng)
meter_comparison = compare_groups(df_converged_draws, "meter", 2, 3,
    ↪ params_to_check, rng)

```

```
df_group_comparisons = pd.concat([condition_comparison, meter_comparison],
    ↪ ignore_index=True)

pd.set_option("display.width", 160)
print(df_group_comparisons.round(3).to_string(index=False))
```

	comparison	parameter	cued_mean	spontaneous_mean	diff_mean	hdi_94_lo
hdi_94_hi	credible	2_mean 3_mean				
cued - spontaneous		tau	27.546	22.827	4.719	3.346
6.005	True	NaN NaN				
cued - spontaneous		s	0.167	0.164	0.003	-0.058
0.063	False	NaN NaN				
cued - spontaneous		sigma_sw	0.428	0.363	0.065	-0.171
0.321	False	NaN NaN				
cued - spontaneous		sigma_0	0.283	0.285	-0.002	-0.029
0.026	False	NaN NaN				
	2 - 3	tau	NaN	NaN	-1.072	-2.425
0.324	False	24.509 25.581				
	2 - 3	s	NaN	NaN	0.006	-0.055
0.064	False	0.168 0.162				
	2 - 3	sigma_sw	NaN	NaN	0.081	-0.163
0.318	False	0.437 0.356				
	2 - 3	sigma_0	NaN	NaN	0.081	0.052
0.109	True	0.327 0.246				

Reading this table: tau differs credibly by condition, but that's not a new finding – it's fully explained by the already-reported cue-to-tapping reaction lag (Section 4) plus the experimenter's own `switch_beat` schedule, not a new property of cued vs. spontaneous switching. `s` and `sigma_sw` show no credible difference either way – switches are equally sharp and equally noisy-at-the-switch regardless of condition or meter, with this amount of data. The one new, credible result is `sigma_0` (baseline noise, unrelated to the switch) being higher in meter-2 trials than meter-3 trials – see `03_analyse.Rmd` for the plot and discussion.

1.9 8. Save results

```
[11]: df_cued_results.to_csv(f"{OUT_DIR}/{SUBJECT}_cued_tau_fit.csv", index=False)
df_spontaneous_results.to_csv(f"{OUT_DIR}/{SUBJECT}_spontaneous_tau_fit.csv",
    ↪ index=False)
df_posterior_draws.to_csv(f"{OUT_DIR}/{SUBJECT}_posterior_draws.csv",
    ↪ index=False)
df_group_comparisons.to_csv(f"{OUT_DIR}/{SUBJECT}_group_comparisons.csv",
    ↪ index=False)

print("Saved:")
print(f" - {OUT_DIR}/{SUBJECT}_cued_tau_fit.csv")
print(f" - {OUT_DIR}/{SUBJECT}_spontaneous_tau_fit.csv")
print(f" - {OUT_DIR}/{SUBJECT}_posterior_draws.csv")
print(f" - {OUT_DIR}/{SUBJECT}_group_comparisons.csv")
```

Saved:

- /Users/robaru/repos/meterswitch/derived/sub02/sub02_cued_tau_fit.csv
- /Users/robaru/repos/meterswitch/derived/sub02/sub02_spontaneous_tau_fit.csv
- /Users/robaru/repos/meterswitch/derived/sub02/sub02_posterior_draws.csv
- /Users/robaru/repos/meterswitch/derived/sub02/sub02_group_comparisons.csv